

Problem Statement:

From an Image dataset we need to find N similar images on a given query image

Importing Data:

Loading Data using Curl-WGet:

```
In [ ]: !wget --header="Host: doc-10-1c-docs.googleusercontent.com" --header="User-Agent: Mozilla/5.0 (Windows NT 10.0; Win64; x64) AppleWebKit/537.36 (KHTML, like Gecko) Chrome/84.0.
```

Unzipping the folder:

```
In [ ]: !unzip dataset.zip
```

- Using CurlWget method takes couple of seconds to load the data.

Note1: When the PC gets turned off all the data you loaded through it will get lost. So you need to load the data again when you run the below code snippets.

Note2: CurlWget gives a .zip file. So we need to unzip the folders before using it.

Mounting Google Drive:

To access/create files in google drive

```
In [ ]: from google.colab import drive  
drive.mount('/content/drive')
```

Loading Dependencies:

```
In [ ]: # For commands  
import os  
os.chdir('/content/')  
import time  
from tqdm.notebook import tqdm  
import warnings  
warnings.filterwarnings('ignore')  
# For array manipulation  
import numpy as np  
import pandas as pd  
import pandas.util.testing as tm  
# For visualization  
import matplotlib.pyplot as plt  
import matplotlib.ticker as ticker  
import seaborn as sns  
import cv2  
import imageio as io  
from pylab import *  
from sklearn.manifold import TSNE  
#For model performance  
from sklearn.metrics import pairwise_distances  
from sklearn.metrics.pairwise import cosine_similarity  
from sklearn.externals import joblib  
#For model training  
from sklearn.model_selection import train_test_split  
from sklearn.cluster import KMeans  
import tensorflow as tf  
from sklearn.neighbors import KNeighborsClassifier  
from tensorflow.keras.models import Sequential  
from keras.applications.vgg16 import preprocess_input  
from tensorflow.keras.layers import Conv2D, Activation, MaxPooling2D, UpSampling2D  
from tensorflow.keras.optimizers import Adam, Adagrad, RMSprop  
from tensorflow.keras.callbacks import EarlyStopping, ModelCheckpoint  
from keras import backend as K  
from keras.models import load_model  
from keras.preprocessing import image
```

- I am using tensorflow.keras (2.0 version) throughout this assignment.
- I am doing this case study in google colab. So all the paths will be specified according to colab directories.

```
In [ ]: file_path = os.listdir('dataset')  
print(len(file_path))
```

4738

- Their are of total 4738 images.
- Since the dataset is small we need to take of model training so that it won't get overfit or underfit.
- Interestingly all the images are in same resolution. (512x512)

Splitting the data into Train & Test sets:

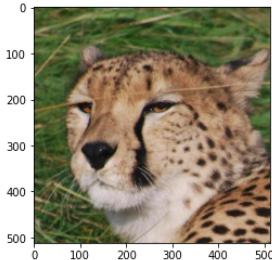
```
In [ ]: train_files, test_files = train_test_split(file_path, test_size = 0.15)  
print(len(train_files))  
print(len(test_files))  
  
train_files = pd.DataFrame(train_files,columns=['filepath'])  
test_files = pd.DataFrame(test_files,columns=['filepath'])  
#converting into .csv file for future reference.  
train_files.to_csv('/content/drive/My Drive/train_file.csv')  
test_files.to_csv('/content/drive/My Drive/test_file.csv')
```

```
In [ ]: #Loading csv files.  
train_files = list(pd.read_csv('/content/drive/My Drive/image_similarity/train_file.csv')['filepath'])  
test_files = list(pd.read_csv('/content/drive/My Drive/image_similarity/test_file.csv')['filepath'])
```

- Taking the file path of all the images and splitting that list into train and test. So that their will be no need of splitting the data again in future and we can access those sets directly by storing it in drive.
- Storing the end result into .csv format so that their will be no data leakage problems in future.

```
In [ ]: img = cv2.imread('/content/dataset/'+train_files[0])
img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
plt.imshow(img)
```

```
Out[3]: <matplotlib.image.AxesImage at 0x7f32798a6748>
```



Reading Images:

```
In [ ]: def image2array(file_array):
    """
    Reading and Converting images into numpy array by taking path of images.
    Arguments:
    file_array - (list) - list of file(path) names
    Returns:
    A numpy array of images. (np.ndarray)
    """

    image_array = []
    for path in tqdm(file_array):
        img = cv2.imread('/content/dataset/' + path)
        img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
        img = cv2.resize(img, (224, 224))
        image_array.append(np.array(img))
    image_array = np.array(image_array)
    image_array = image_array.reshape(image_array.shape[0], 224, 224, 3)
    image_array = image_array.astype('float32')
    image_array /= 255
    return np.array(image_array)
```

```
In [ ]: train_data = image2array(train_files)
print("Length of training dataset:", train_data.shape)
test_data = image2array(test_files)
print("Length of test dataset:", test_data.shape)
```

```
HBox(children=(FloatProgress(value=0.0, max=4027.0), HTML(value='')))
```

```
Length of training dataset: (4027, 224, 224, 3)
```

```
HBox(children=(FloatProgress(value=0.0, max=711.0), HTML(value='')))
```

```
Length of test dataset: (711, 224, 224, 3)
```

Model Architecture & Model Training:

```
In [ ]: def encoder_decoder_model():
    """
    Used to build Convolutional Autoencoder model architecture to get compressed image data which is easier to process.
    Returns:
    Auto encoder model
    """

    #Encoder
    model = Sequential(name='Convolutional_AutoEncoder_Model')
    model.add(Conv2D(64, kernel_size=(3, 3), activation='relu', input_shape=(224, 224, 3), padding='same', name='Encoding_Conv2D_1'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='same', name='Encoding_MaxPooling2D_1'))
    model.add(Conv2D(128, kernel_size=(3, 3), strides=1, kernel_regularizer = tf.keras.regularizers.L2(0.001), activation='relu', padding='same', name='Encoding_Conv2D_2'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='same', name='Encoding_MaxPooling2D_2'))
    model.add(Conv2D(256, kernel_size=(3, 3), activation='relu', kernel_regularizer = tf.keras.regularizers.L2(0.001), padding='same', name='Encoding_Conv2D_3'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='same', name='Encoding_MaxPooling2D_3'))
    model.add(Conv2D(512, kernel_size=(3, 3), activation='relu', kernel_regularizer = tf.keras.regularizers.L2(0.001), padding='same', name='Encoding_Conv2D_4'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='valid', name='Encoding_MaxPooling2D_4'))
    model.add(Conv2D(512, kernel_size=(3, 3), activation='relu', padding='same', name='Encoding_Conv2D_5'))
    model.add(MaxPooling2D(pool_size=(2, 2), strides=2, padding='valid', name='Encoding_MaxPooling2D_5'))

    #Decoder
    model.add(Conv2D(512, kernel_size=(3, 3), kernel_regularizer = tf.keras.regularizers.L2(0.001), activation='relu', padding='same', name='Decoding_Conv2D_1'))
    model.add(UpSampling2D((2, 2), name='Decoding_Upsampling2D_1'))
    model.add(Conv2D(512, kernel_size=(3, 3), kernel_regularizer = tf.keras.regularizers.L2(0.001), activation='relu', padding='same', name='Decoding_Conv2D_2'))
    model.add(UpSampling2D((2, 2), name='Decoding_Upsampling2D_2'))
    model.add(Conv2D(256, kernel_size=(3, 3), kernel_regularizer = tf.keras.regularizers.L2(0.001), activation='relu', padding='same', name='Decoding_Conv2D_3'))
    model.add(UpSampling2D((2, 2), name='Decoding_Upsampling2D_3'))
    model.add(Conv2D(128, kernel_size=(3, 3), activation='relu', kernel_regularizer = tf.keras.regularizers.L2(0.001), padding='same', name='Decoding_Conv2D_4'))
    model.add(UpSampling2D((2, 2), name='Decoding_Upsampling2D_4'))
    model.add(Conv2D(64, kernel_size=(3, 3), activation='relu', kernel_regularizer = tf.keras.regularizers.L2(0.001), padding='same', name='Decoding_Conv2D_5'))
    model.add(UpSampling2D((2, 2), name='Decoding_Upsampling2D_5'))
    model.add(Conv2D(3, kernel_size=(3, 3), padding='same', activation='sigmoid', name='Decoding_Output'))
    return model
```

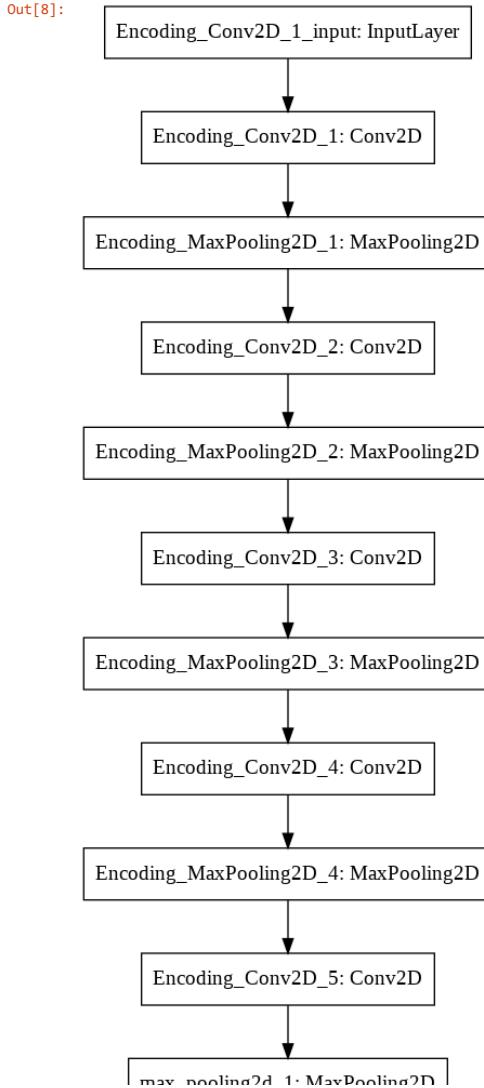
```
In [ ]: model = encoder_decoder_model()
model.summary()
print("\n")
tf.keras.utils.plot_model(model, to_file='/content/drive/My Drive/model.png')

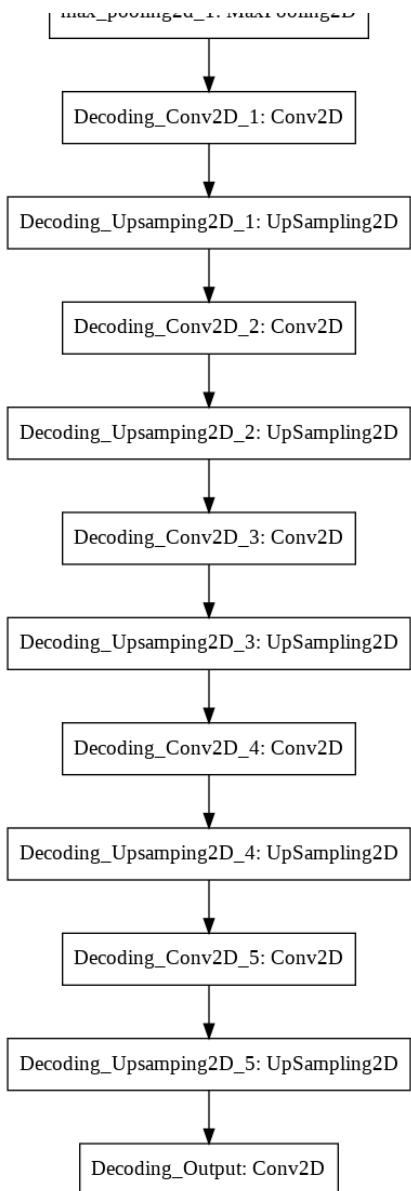
Model: "Convolutional_AutoEncoder_Model"

```

Layer (type)	Output Shape	Param #
Encoding_Conv2D_1 (Conv2D)	(None, 224, 224, 64)	1792
Encoding_MaxPooling2D_1 (Max)	(None, 112, 112, 64)	0
Encoding_Conv2D_2 (Conv2D)	(None, 112, 112, 128)	73856
Encoding_MaxPooling2D_2 (Max)	(None, 56, 56, 128)	0
Encoding_Conv2D_3 (Conv2D)	(None, 56, 56, 256)	295168
Encoding_MaxPooling2D_3 (Max)	(None, 28, 28, 256)	0
Encoding_Conv2D_4 (Conv2D)	(None, 28, 28, 512)	1180160
Encoding_MaxPooling2D_4 (Max)	(None, 14, 14, 512)	0
Encoding_Conv2D_5 (Conv2D)	(None, 14, 14, 512)	2359808
max_pooling2d_1 (MaxPooling2D)	(None, 7, 7, 512)	0
Decoding_Conv2D_1 (Conv2D)	(None, 7, 7, 512)	2359808
Decoding_Upsampling2D_1 (UpSampling2D)	(None, 14, 14, 512)	0
Decoding_Conv2D_2 (Conv2D)	(None, 14, 14, 512)	2359808
Decoding_Upsampling2D_2 (UpSampling2D)	(None, 28, 28, 512)	0
Decoding_Conv2D_3 (Conv2D)	(None, 28, 28, 256)	1179904
Decoding_Upsampling2D_3 (UpSampling2D)	(None, 56, 56, 256)	0
Decoding_Conv2D_4 (Conv2D)	(None, 56, 56, 128)	295040
Decoding_Upsampling2D_4 (UpSampling2D)	(None, 112, 112, 128)	0
Decoding_Conv2D_5 (Conv2D)	(None, 112, 112, 64)	73792
Decoding_Upsampling2D_5 (UpSampling2D)	(None, 224, 224, 64)	0
Decoding_Output (Conv2D)	(None, 224, 224, 3)	1731

Total params: 10,180,867
Trainable params: 10,180,867
Non-trainable params: 0





Hyper parameter Tuning:

```
In [ ]: parameters = {'Adagrad':[0.01,0.001,0.0001,0.00001], 'Adam':[0.01,0.001,0.0001,0.00001], 'Rmsprop':[0.01,0.001,0.0001,0.00001]}
result = []
for i in parameters.keys():
    print("{} as an optimizer:".format(i))
    values = parameters[i]
    result_ = []
    for learning_rate in values:
        print("\t\tUsing learning_rate: {}".format(learning_rate))
        model = encoder_decoder_model()
        if i=='Adam':
            optimizer = Adam(learning_rate=learning_rate)
        elif i=="Adagrad":
            optimizer = Adagrad(learning_rate=learning_rate)
        else:
            optimizer = RMSprop(learning_rate=learning_rate)
        model.compile(optimizer=optimizer, loss='mse') # compiling
        model.fit(train_data, train_data, epochs=5, batch_size=32, validation_data=(test_data,test_data)) # fitting data
        result_.append(model.history.history) # taking result to judge the best parameters.
    print()
result.append(result_)

Adagrad as an optimizer:
    Using learning_rate: 0.01
Epoch 1/5
126/126 [=====] - 20s 158ms/step - loss: 0.0533 - val_loss: 0.0512
Epoch 2/5
126/126 [=====] - 19s 148ms/step - loss: 0.0504 - val_loss: 0.0487
Epoch 3/5
126/126 [=====] - 19s 149ms/step - loss: 0.0475 - val_loss: 0.0453
Epoch 4/5
126/126 [=====] - 19s 148ms/step - loss: 0.0431 - val_loss: 0.0401
Epoch 5/5
126/126 [=====] - 19s 148ms/step - loss: 0.0369 - val_loss: 0.0333
    Using learning_rate: 0.001
Epoch 1/5
126/126 [=====] - 19s 149ms/step - loss: 0.0562 - val_loss: 0.0555
Epoch 2/5
126/126 [=====] - 19s 148ms/step - loss: 0.0558 - val_loss: 0.0552
Epoch 3/5
126/126 [=====] - 19s 147ms/step - loss: 0.0555 - val_loss: 0.0549
Epoch 4/5
126/126 [=====] - 19s 148ms/step - loss: 0.0552 - val_loss: 0.0546
Epoch 5/5
126/126 [=====] - 19s 147ms/step - loss: 0.0549 - val_loss: 0.0544
    Using learning_rate: 0.0001
Epoch 1/5
126/126 [=====] - 19s 149ms/step - loss: 0.0568 - val_loss: 0.0563
Epoch 2/5
126/126 [=====] - 19s 147ms/step - loss: 0.0567 - val_loss: 0.0562
Epoch 3/5
126/126 [=====] - 19s 148ms/step - loss: 0.0567 - val_loss: 0.0562
Epoch 4/5
126/126 [=====] - 19s 148ms/step - loss: 0.0567 - val_loss: 0.0562
Epoch 5/5
126/126 [=====] - 19s 148ms/step - loss: 0.0566 - val_loss: 0.0561
    Using learning_rate: 1e-05
Epoch 1/5
126/126 [=====] - 19s 148ms/step - loss: 0.0562 - val_loss: 0.0557
Epoch 2/5
126/126 [=====] - 19s 148ms/step - loss: 0.0562 - val_loss: 0.0557
Epoch 3/5
126/126 [=====] - 19s 148ms/step - loss: 0.0562 - val_loss: 0.0557
Epoch 4/5
126/126 [=====] - 19s 148ms/step - loss: 0.0562 - val_loss: 0.0557
Epoch 5/5
126/126 [=====] - 19s 148ms/step - loss: 0.0562 - val_loss: 0.0557

Adam as an optimizer:
    Using learning_rate: 0.01
Epoch 1/5
126/126 [=====] - 18s 146ms/step - loss: 0.0225 - val_loss: 0.0145
Epoch 2/5
126/126 [=====] - 18s 144ms/step - loss: 0.0131 - val_loss: 0.0147
Epoch 3/5
126/126 [=====] - 18s 144ms/step - loss: 0.0120 - val_loss: 0.0112
Epoch 4/5
126/126 [=====] - 18s 145ms/step - loss: 0.0117 - val_loss: 0.0114
Epoch 5/5
126/126 [=====] - 18s 145ms/step - loss: 0.0098 - val_loss: 0.0090
    Using learning_rate: 0.001
Epoch 1/5
126/126 [=====] - 19s 149ms/step - loss: 0.0287 - val_loss: 0.0147
Epoch 2/5
126/126 [=====] - 19s 148ms/step - loss: 0.0123 - val_loss: 0.0109
Epoch 3/5
126/126 [=====] - 19s 148ms/step - loss: 0.0104 - val_loss: 0.0104
Epoch 4/5
126/126 [=====] - 19s 149ms/step - loss: 0.0098 - val_loss: 0.0095
Epoch 5/5
126/126 [=====] - 19s 148ms/step - loss: 0.0093 - val_loss: 0.0093
    Using learning_rate: 0.0001
Epoch 1/5
126/126 [=====] - 19s 148ms/step - loss: 0.0485 - val_loss: 0.0306
Epoch 2/5
126/126 [=====] - 19s 147ms/step - loss: 0.0221 - val_loss: 0.0179
Epoch 3/5
126/126 [=====] - 19s 148ms/step - loss: 0.0155 - val_loss: 0.0140
Epoch 4/5
126/126 [=====] - 19s 148ms/step - loss: 0.0130 - val_loss: 0.0125
Epoch 5/5
126/126 [=====] - 19s 147ms/step - loss: 0.0120 - val_loss: 0.0117
    Using learning_rate: 1e-05
Epoch 1/5
126/126 [=====] - 19s 149ms/step - loss: 0.0559 - val_loss: 0.0548
Epoch 2/5
126/126 [=====] - 18s 147ms/step - loss: 0.0546 - val_loss: 0.0534
Epoch 3/5
126/126 [=====] - 19s 147ms/step - loss: 0.0534 - val_loss: 0.0525
Epoch 4/5
126/126 [=====] - 19s 147ms/step - loss: 0.0526 - val_loss: 0.0517
Epoch 5/5
126/126 [=====] - 19s 147ms/step - loss: 0.0517 - val_loss: 0.0506

Rmsprop as an optimizer:
    Using learning_rate: 0.01
Epoch 1/5
126/126 [=====] - 19s 147ms/step - loss: 0.0536 - val_loss: 0.0511
```

```

Epoch 2/5
126/126 [=====] - 18s 145ms/step - loss: 0.0515 - val_loss: 0.0511
Epoch 3/5
126/126 [=====] - 18s 145ms/step - loss: 0.0515 - val_loss: 0.0511
Epoch 4/5
126/126 [=====] - 18s 145ms/step - loss: 0.0515 - val_loss: 0.0511
Epoch 5/5
126/126 [=====] - 18s 145ms/step - loss: 0.0514 - val_loss: 0.0511
    Using learning_rate: 0.001
Epoch 1/5
126/126 [=====] - 19s 150ms/step - loss: 0.0279 - val_loss: 0.0184
Epoch 2/5
126/126 [=====] - 19s 149ms/step - loss: 0.0179 - val_loss: 0.0160
Epoch 3/5
126/126 [=====] - 19s 149ms/step - loss: 0.0157 - val_loss: 0.0124
Epoch 4/5
126/126 [=====] - 19s 149ms/step - loss: 0.0143 - val_loss: 0.0129
Epoch 5/5
126/126 [=====] - 19s 148ms/step - loss: 0.0133 - val_loss: 0.0129
    Using learning_rate: 0.0001
Epoch 1/5
126/126 [=====] - 19s 150ms/step - loss: 0.0446 - val_loss: 0.0281
Epoch 2/5
126/126 [=====] - 19s 149ms/step - loss: 0.0229 - val_loss: 0.0202
Epoch 3/5
126/126 [=====] - 19s 149ms/step - loss: 0.0190 - val_loss: 0.0188
Epoch 4/5
126/126 [=====] - 19s 149ms/step - loss: 0.0175 - val_loss: 0.0174
Epoch 5/5
126/126 [=====] - 19s 149ms/step - loss: 0.0165 - val_loss: 0.0157
    Using learning_rate: 1e-05
Epoch 1/5
126/126 [=====] - 19s 150ms/step - loss: 0.0552 - val_loss: 0.0541
Epoch 2/5
126/126 [=====] - 19s 148ms/step - loss: 0.0540 - val_loss: 0.0531
Epoch 3/5
126/126 [=====] - 19s 148ms/step - loss: 0.0533 - val_loss: 0.0525
Epoch 4/5
126/126 [=====] - 19s 149ms/step - loss: 0.0527 - val_loss: 0.0519
Epoch 5/5
126/126 [=====] - 19s 149ms/step - loss: 0.0520 - val_loss: 0.0511

```

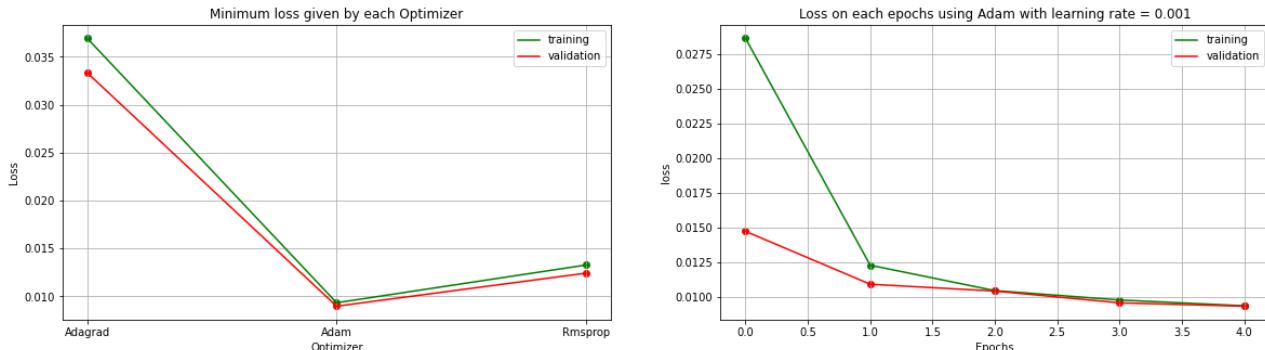
```

In [ ]: def plot_(x,y1,y2,row,col,ind,title,xlabel,ylabel,label,isimage=False,color='r'):

    """
    This function is used for plotting images and graphs (Visualization of end results of model training)
    Arguments:
    x - (np.ndarray or list) - an image array
    y1 - (list) - for plotting graph on left side.
    y2 - (list) - for plotting graph on right side.
    row - (int) - row number of subplot
    col - (int) - column number of subplot
    ind - (int) - index number of subplot
    title - (string) - title of the plot
    xlabel - (list) - labels of x axis
    ylabel - (list) - labels of y axis
    label - (string) - for adding legend in the plot
    isimage - (boolean) - True in case of image else False
    color - (char) - color of the plot (prefered green for training and red for testing).
    """
    plt.subplot(row,col,ind)
    if isimage:
        plt.imshow(x)
        plt.title(title)
        plt.axis('off')
    else:
        plt.plot(y1,label=label,color='g'); plt.scatter(x,y1,color='g')
        if y2 != '': plt.plot(y2,color=color,label='validation'); plt.scatter(x,y2,color=color)
        plt.grid()
        plt.legend()
        plt.title(title); plt.xlabel(xlabel); plt.ylabel(ylabel)

In [ ]: min_train = []; min_val = []
rates = list(parameters.keys())
epochs = [0,1,2,3,4]
for i in result:
    train = []; val = []
    for j in i:
        train.append(min(j['loss'])); val.append(min(j['val_loss'])) # taking minimum loss of each optimizer over all Learning rates.
    min_train.append(min(train)); min_val.append(min(val))
plt.figure(figsize=(20,5))
plot_(rates,min_train,min_val,1,2,1,'Minimum loss given by each Optimizer','Optimizer','Loss','training',False,'r')
# plotting the result of adam with Learning rate = 0.001 .
plot_(epochs,result[1][1]['loss'],result[1][1]['val_loss'],1,2,2,'Loss on each epochs using Adam with learning rate = 0.001','Epochs','loss','training',False,'r')
plt.show()

```



- After training the model with different optimizers(Adagrad, Adam, Rmsprop), adam giving the least local minimum loss on training with learning rate = 0.001.
- Both training loss and validation loss are almost equal and we can see our model training is not prone to overfitting and underfitting.
- Achieved 0.0093 on five epochs using Adam(0.001)(optimal value).

Training the model with Best Optimizer with Best Learning Rate:

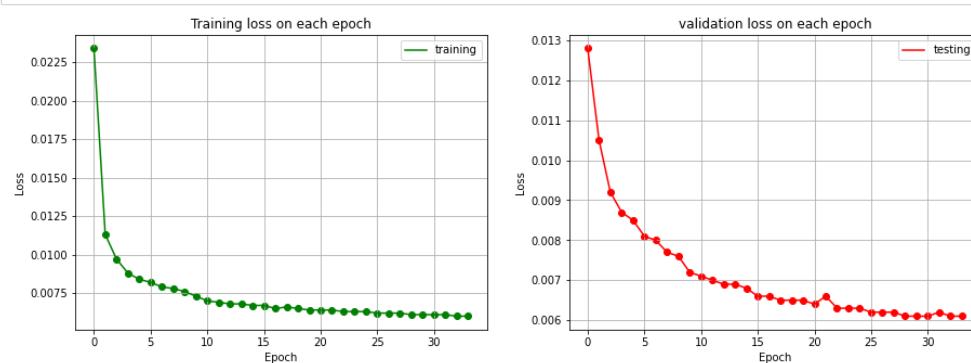
```
In [ ]: optimizer = Adam(learning_rate=0.001)
model = encoder_decoder_model()
model.compile(optimizer=optimizer, loss='mse')
early_stopping = EarlyStopping(monitor='val_loss', mode='min', verbose=1, patience=6, min_delta=0.0001)
checkpoint = ModelCheckpoint('/content/drive/My Drive/encoder_model.h5', monitor='val_loss', mode='min', save_best_only=True)
model.fit(train_data, train_data, epochs=35, batch_size=32, validation_data=(test_data, test_data), callbacks=[early_stopping, checkpoint])
```

Epoch 1/35
126/126 [=====] - 20s 156ms/step - loss: 0.0234 - val_loss: 0.0128
Epoch 2/35
126/126 [=====] - 20s 156ms/step - loss: 0.0113 - val_loss: 0.0105
Epoch 3/35
126/126 [=====] - 20s 157ms/step - loss: 0.0097 - val_loss: 0.0092
Epoch 4/35
126/126 [=====] - 20s 158ms/step - loss: 0.0088 - val_loss: 0.0087
Epoch 5/35
126/126 [=====] - 20s 157ms/step - loss: 0.0084 - val_loss: 0.0085
Epoch 6/35
126/126 [=====] - 20s 157ms/step - loss: 0.0082 - val_loss: 0.0081
Epoch 7/35
126/126 [=====] - 20s 157ms/step - loss: 0.0079 - val_loss: 0.0080
Epoch 8/35
126/126 [=====] - 20s 155ms/step - loss: 0.0078 - val_loss: 0.0077
Epoch 9/35
126/126 [=====] - 20s 157ms/step - loss: 0.0076 - val_loss: 0.0076
Epoch 10/35
126/126 [=====] - 20s 155ms/step - loss: 0.0073 - val_loss: 0.0072
Epoch 11/35
126/126 [=====] - 19s 155ms/step - loss: 0.0072 - val_loss: 0.0071
Epoch 12/35
126/126 [=====] - 19s 154ms/step - loss: 0.0070 - val_loss: 0.0070
Epoch 13/35
126/126 [=====] - 19s 155ms/step - loss: 0.0069 - val_loss: 0.0069
Epoch 14/35
126/126 [=====] - 19s 155ms/step - loss: 0.0068 - val_loss: 0.0069
Epoch 15/35
126/126 [=====] - 19s 155ms/step - loss: 0.0068 - val_loss: 0.0068
Epoch 16/35
126/126 [=====] - 19s 155ms/step - loss: 0.0067 - val_loss: 0.0067
Epoch 17/35
126/126 [=====] - 19s 155ms/step - loss: 0.0067 - val_loss: 0.0066
Epoch 18/35
126/126 [=====] - 20s 156ms/step - loss: 0.0065 - val_loss: 0.0066
Epoch 19/35
126/126 [=====] - 20s 155ms/step - loss: 0.0066 - val_loss: 0.0065
Epoch 20/35
126/126 [=====] - 19s 154ms/step - loss: 0.0065 - val_loss: 0.0065
Epoch 21/35
126/126 [=====] - 20s 156ms/step - loss: 0.0064 - val_loss: 0.0065
Epoch 22/35
126/126 [=====] - 20s 155ms/step - loss: 0.0064 - val_loss: 0.0064
Epoch 23/35
126/126 [=====] - 19s 147ms/step - loss: 0.0064 - val_loss: 0.0066
Epoch 24/35
126/126 [=====] - 20s 155ms/step - loss: 0.0063 - val_loss: 0.0063
Epoch 25/35
126/126 [=====] - 20s 155ms/step - loss: 0.0063 - val_loss: 0.0063
Epoch 26/35
126/126 [=====] - 19s 147ms/step - loss: 0.0063 - val_loss: 0.0063
Epoch 27/35
126/126 [=====] - 20s 155ms/step - loss: 0.0062 - val_loss: 0.0062
Epoch 28/35
126/126 [=====] - 20s 158ms/step - loss: 0.0062 - val_loss: 0.0062
Epoch 29/35
126/126 [=====] - 20s 155ms/step - loss: 0.0062 - val_loss: 0.0062
Epoch 30/35
126/126 [=====] - 20s 155ms/step - loss: 0.0061 - val_loss: 0.0061
Epoch 31/35
126/126 [=====] - 20s 155ms/step - loss: 0.0061 - val_loss: 0.0061
Epoch 32/35
126/126 [=====] - 20s 156ms/step - loss: 0.0061 - val_loss: 0.0061
Epoch 33/35
126/126 [=====] - 18s 147ms/step - loss: 0.0061 - val_loss: 0.0062
Epoch 34/35
126/126 [=====] - 20s 156ms/step - loss: 0.0060 - val_loss: 0.0061
Epoch 35/35
126/126 [=====] - 19s 154ms/step - loss: 0.0060 - val_loss: 0.0061

Out[20]: <tensorflow.python.keras.callbacks.History at 0x7f2b0cf774a8>

Plotting loss on each epoch:

```
In [ ]: #model.history.history
plt.figure(figsize=(15,5))
epochs = [i for i in range(34)]
plot_(epochs,loss,'',1,2,1,'Training loss on each epoch','Epoch','Loss','training',False,'g')
plot_(epochs,val_loss,'',1,2,2,'validation loss on each epoch','Epoch','Loss','testing',False,'r')
```



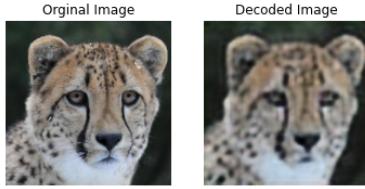
- There is a strict decrement in the loss for both training and validation.
- After 30 epochs achieved 0.0061 loss

Loading the trained model:

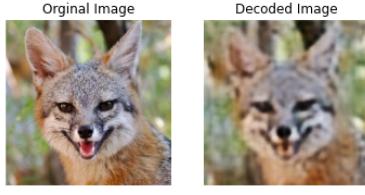
```
In [ ]: model = load_model("/content/drive/My Drive/encoder_model.h5")
model.compile(optimizer=optimizer, loss='mse')
```

Model Testing:

```
In [ ]: sample_image = train_data[7]
sample_image = np.expand_dims(sample_image, axis=0)
image = model.predict(sample_image)
plot_(sample_image[0,:,:,:],'',1,2,1,"Orginal Image","","","",True)
plot_(image[0,:,:,:],'',1,2,2,"Decoded Image","","","",True)
plt.show()
```

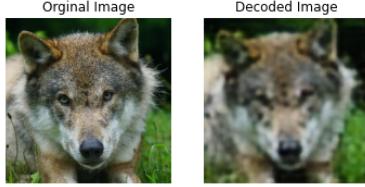


```
In [ ]: sample_image = train_data[2396]
sample_image = np.expand_dims(sample_image, axis=0)
image = model.predict(sample_image)
plot_(sample_image[0,:,:,:],'',1,2,1,"Orginal Image","","","",True)
plot_(image[0,:,:,:],'',1,2,2,"Decoded Image","","","",True)
plt.show()
```



Restoring the Best Model using Model Checkpoint:

```
In [ ]: model = load_model('/content/drive/My Drive/encoder_model.h5')
sample_image = train_data[15]
sample_image = np.expand_dims(sample_image, axis=0)
image = model.predict(sample_image)
plot_(train_data[15],'',1,2,1,"Orginal Image","","","",True)
plot_(image[0,:,:,:],'',1,2,2,"Decoded Image","","","",True)
plt.show()
```



- Restorations seems really satisfactory. Images on the left side are original images whereas images on the right side are restored from compressed representation.
- Decoded image is much flexible and efficient to work rather than working with original image since the compressed representation takes 8 times less space to original image.

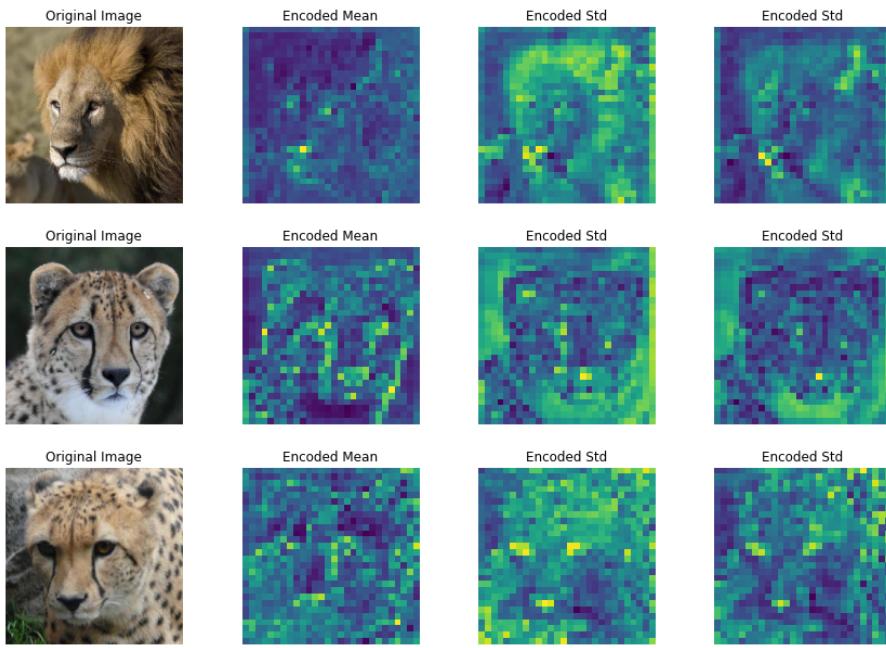
Feature Extraction:

```
In [ ]: from keras import backend as K
def feature_extraction(model, data, layer = 4):

    """
    Creating a function to run the initial layers of the encoder model. (to get feature extraction from any layer of the model)
    Arguments:
    model - (Auto encoder model) - Trained model
    data - (np.ndarray) - list of images to get feature extraction from trained model
    layer - (int) - from which layer to take the features(by default = 4)
    Returns:
    pooled_array - (np.ndarray) - array of extracted features of given images
    """

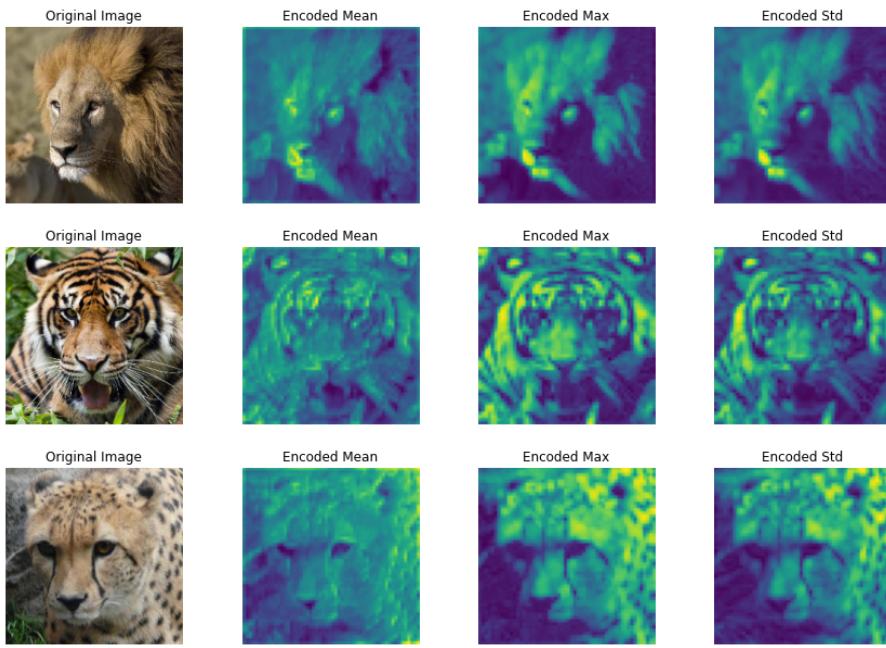
    encoded = K.function([model.layers[0].input],[model.layers[layer].output])
    encoded_array = encoded([data])[0]
    pooled_array = encoded_array.max(axis=-1)
    return pooled_array
encoded = feature_extraction(model,train_data[:10],12)
```

```
In [ ]: for index in [2,7,9]: # 3 random images
    plt.figure(figsize=(15,3))
    plot_(train_data[index],'',1,4,1,"Original Image","",",",True)
    plot_(encoded[index].mean(axis=-1),'',1,4,2,"Encoded Mean","",",",True)
    plot_(encoded[index].max(axis=-1),'',1,4,3,"Encoded Max","",",",True)
    plot_(encoded[index].std(axis=-1),'',1,4,4,"Encoded Std","",",",True)
plt.show()
```



- Extracting features from the 12th layer.
- Dark pixels(yellow) indicates the high activation which helps in differentiating with other images.

```
In [ ]: encoded = feature_extraction(model,train_data[:10],9)
for index in [2,6,9]: # 3 random images
    plt.figure(figsize=(15,3))
    plot_(train_data[index],'',1,4,1,"Original Image","",",",True)
    plot_(encoded[index].mean(axis=-1),'',1,4,2,"Encoded Mean","",",",True)
    plot_(encoded[index].max(axis=-1),'',1,4,3,"Encoded Max","",",",True)
    plot_(encoded[index].std(axis=-1),'',1,4,4,"Encoded Std","",",",True)
plt.show()
```



Extracting features from the 9th layer. Dark pixels(yellow) indicates the high activation which helps in differentiating with other images.

1. High activation on mane of lion.
2. Nose and lines on tigers.
3. Dots and lines on the nose for cheetah.
4. Nose on foxes. These activation helps in label classification.

```
In [ ]: def get_batches(data, batch_size=1000):
    """
    Taking batch of images for extraction of images.
    Arguments:
    data - (np.ndarray or list) - list of image array to get extracted features.
    batch_size - (int) - Number of images per each batch
    Returns:
    list - extracted features of each images
    """

    if len(data) < batch_size:
        return [data]
    n_batches = len(data) // batch_size

    # If batches fit exactly into the size of df.
    if len(data) % batch_size == 0:
        return [data[i*batch_size:(i+1)*batch_size] for i in range(n_batches)]

    # If there is a remainder.
    else:
        return [data[i*batch_size:min((i+1)*batch_size, len(data))] for i in range(n_batches+1)]
```

```
In [ ]: d = np.concatenate([train_data,test_data],axis=0)
d.shape
```

```
Out[25]: (4738, 224, 224, 3)
```

```
In [ ]: X_encoded = []
i=0
# Iterate through the full training set.
for batch in get_batches(d, batch_size=300):
    i+=1
    # This Line runs our pooling function on the model for each batch.
    X_encoded.append(feature_extraction(model, batch),12)

X_encoded = np.concatenate(X_encoded)
```

```
In [ ]: X_encoded.shape
```

```
Out[27]: (4738, 7, 7, 512)
```

```
In [ ]: X_encoded_reshape = X_encoded.reshape(X_encoded.shape[0], X_encoded.shape[1]*X_encoded.shape[2]*X_encoded.shape[3])
print('Encoded shape:', X_encoded_reshape.shape)
np.save('/content/drive/My Drive/X_encoded_compressed.npy',X_encoded_reshape)

Encoded shape: (4738, 25088)
```

```
In [ ]: X_encoded = np.load('/content/drive/My Drive/X_encoded_compressed.npy')
X_encoded.shape
```

```
Out[84]: (4738, 25088)
```

```
In [ ]: lisps=train_files
lisps.extend(test_files)
print(len(lisps))

4738
```

Dimensionality Reduction through T-SNE:

```
In [ ]: transform = TSNE
trans = transform(n_components=2)
values = trans.fit_transform(X_encoded_reshape)
```

```
In [ ]: def plot_(x,y1,y2,row,col,ind,title,xlabel,ylabel,label,isimage=False,color='b'):

    """
    This function is used for plotting images and graphs (Visualization of end results of model training)
    Arguments:
    x - (np.ndarray or list) - an image array
    y1 - (list) - for plotting graph on left side.
    y2 - (list) - for plotting graph on right side.
    row - (int) - row number of subplot
    col - (int) - column number of subplot
    ind - (int) - index number of subplot
    title - (string) - title of the plot
    xlabel - (list) - labels of x axis
    ylabel - (list) - labels of y axis
    label - (string) - for adding legend in the plot
    isimage - (boolean) - True in case of image else False
    color - (char) - color of the plot (prefered green for training and red for testing).
    """

    plt.subplot(row,col,ind)
    if isimage:
        plt.imshow(x)
        plt.title(title)
        plt.axis('off')
    else:
        plt.plot(y1,label=label,color='g'); plt.scatter(x,y1,color='g')
        if y2=='': plt.plot(y2,color=color,label='validation'); plt.scatter(x,y2,color=color)
        plt.grid()
        plt.legend()
        plt.title(title); plt.xlabel(xlabel); plt.ylabel(ylabel)
```

```
In [ ]: lisps=train_files
lisps.extend(test_files)
print(len(lisps))
```

Clustering Image Data:(K-Means)

Plotting few images of each cluster:

```
In [ ]: K = [4,5,6,7]
for k in K:
    print("if Number of clusters: "+str(k))
    kmeans = KMeans(n_clusters = k, random_state=0).fit(X_encoded_reshape)
    labels=kmeans.labels_
    centroids = kmeans.cluster_centers_
    plt.figure(figsize=(10,5))
    plt.subplot(1,1,1)
    plt.scatter(values[:,0], values[:,1], c= kmeans.labels_.astype(float), s=50, alpha=0.5)
    plt.scatter(centroids[:, 0], centroids[:, 1], c=None, s=50)
    plt.show()
    for row in range(k):
        iter=0
        plt.figure(figsize=(13,3))
        for i,iterator in enumerate(labels):
            if iterator == row:
                img = cv2.imread("/content/dataset/"+lisp[i])
                img = cv2.cvtColor(img, cv2.COLOR_BGR2RGB)
                plot_(img,"","",1,6,iter+1,"cluster="+str(row),"","","",True)
                iter+=1
            if iter>5: break
        plt.show()
print()
```

Output hidden; open in <https://colab.research.google.com> (<https://colab.research.google.com>) to view.

with cluster = 4

1. lions into first cluster.
2. cheetahs into second cluster.
3. foxes into third cluster.
4. Tigers into fourth cluster.

with cluster = 5

1. foxes into first cluster.
2. cheetahs into second cluster.
3. lions into third cluster.
4. Tigers into fourth cluster.
5. Tiger images with different lightning conditions and white tigers into fifth cluster.

with cluster = 6

1. cheetahs into first cluster.
2. lions into second cluster.
3. snow dogs into third, tigers into fourth cluster.
4. leopards into fifth cluster.
5. foxes into sixth cluster. (**optimal parameter**)

with cluster = 7

1. Taking white tigers,cheetahs into separate cluster.

```
In [ ]: #Training the model with optimial K value (6 in our case)
kmeans = KMeans(n_clusters = 6, random_state=0).fit(X_encoded_reshape)
labels=kmeans.labels_
centroids = kmeans.cluster_centers_
```

Storing the model for future reference:

```
In [ ]: kmeans_file = '/content/drive/My Drive/kmeans_model.pkl'
joblib.dump(kmeans,kmeans_file)
```

```
In [ ]: clusters_features = []
cluster_files=[]
for i in [0,1,2,3,4,5]:
    i_cluster = []
    i_labels=[]
    for iter,j in enumerate(kmeans.labels_):
        if j==i:
            i_cluster.append(X_encoded_reshape[iter])
            i_labels.append(lisp[iter])
    i_cluster = np.array(i_cluster)
    clusters_features.append(i_cluster)
    cluster_files.append(i_labels)
```

```
In [ ]: labels=[]
data=[]
files=[]
for iter,i in enumerate(clusters_features):
    data.extend(i)
    labels.extend([iter for i in range(i.shape[0])])
    files.extend(cluster_files[iter])
print(np.array(labels).shape)
print(np.array(data).shape)
print(np.array(files).shape)

(4738,)
(4738, 25088)
(4738,)
```

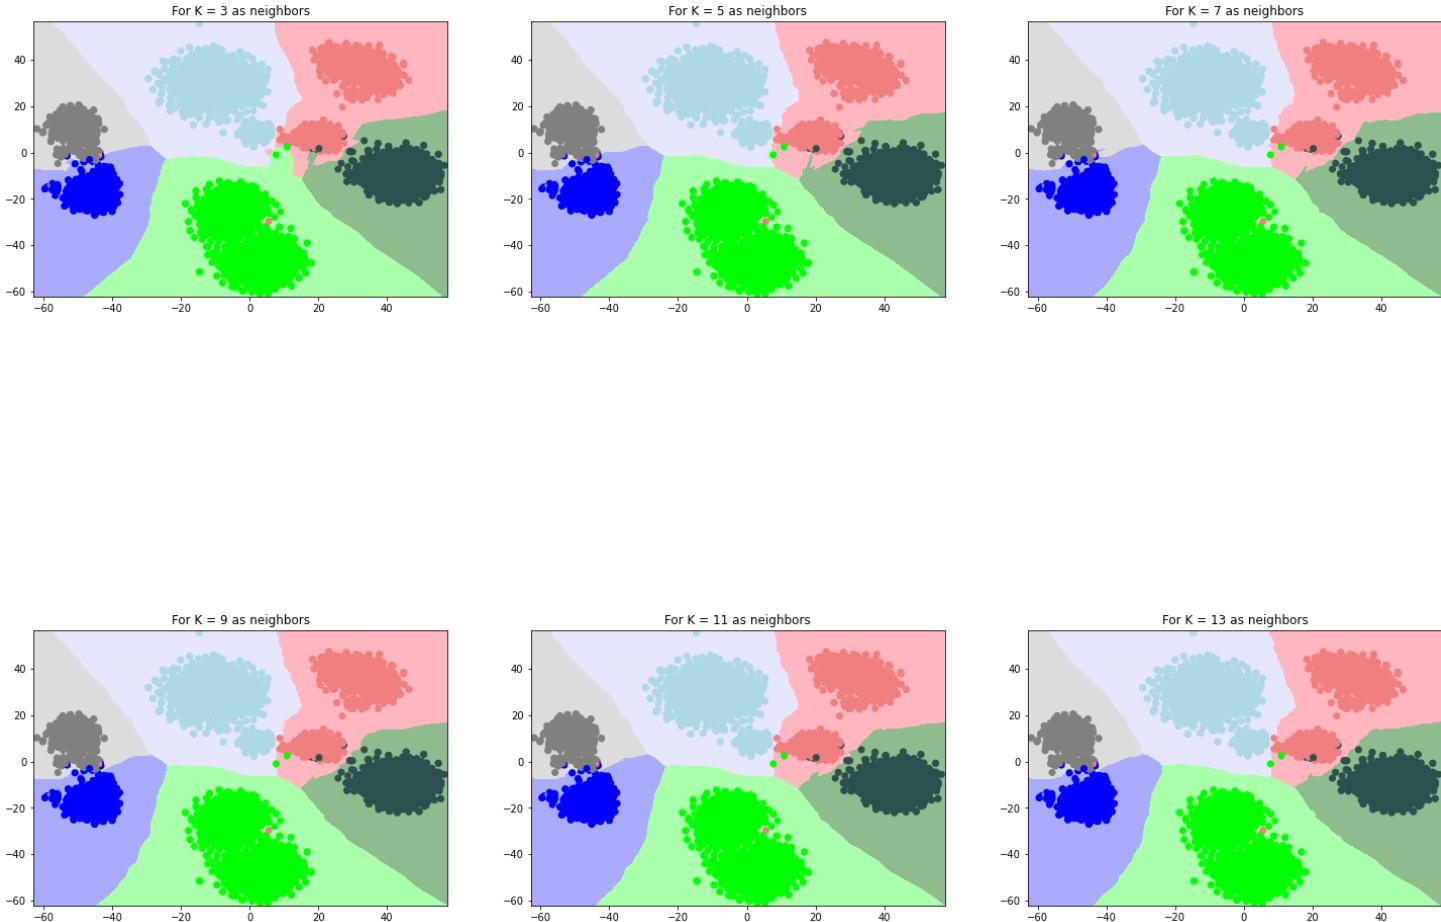
Finding Nearest Neighbors(K-NN):

```
In [ ]: for i in [[3,5,7],[9,11,13]]:
    plt.figure(figsize=(25,5))
    for iter,j in enumerate(i):
        n_neighbors = j
        X = values
        y = labels
        h = .09 # step size in the mesh
        cmap_light = ListedColormap(['#FFB6C1', '#AAFFAA', '#AAAAFF', '#E6E6FA', '#8FBC8F', '#DCDCDC'])
        cmap_bold = ListedColormap(['#F08080', '#00FF00', '#0000FF', '#ADD8E6', '#2F4F4F', '#808080'])
        clf = KNeighborsClassifier(n_neighbors)
        clf.fit(X, y)

        x_min, x_max = X[:, 0].min() - 1, X[:, 0].max() + 1
        y_min, y_max = X[:, 1].min() - 1, X[:, 1].max() + 1
        xx, yy = np.meshgrid(np.arange(x_min, x_max, h), np.arange(y_min, y_max, h))
        Z = clf.predict(xx.ravel(), yy.ravel())

        plt.subplot(1,3,iter+1)
        Z = Z.reshape(xx.shape)
        plt.pcolormesh(xx, yy, Z, cmap=cmap_light)

        # Plot also the training points
        plt.scatter(X[:, 0], X[:, 1], c=y, cmap=cmap_bold)
        plt.xlim(xx.min(), xx.max())
        plt.ylim(yy.min(), yy.max())
        plt.title("For K = {} as neighbors".format(j))
    plt.show()
```



Training the model with optimal hyperparameter:

```
In [ ]: knn = KNeighborsClassifier(n_neighbors=9,algorithm='ball_tree',n_jobs=-1)
knn.fit(np.array(data),np.array(labels))
```

```
Out[21]: KNeighborsClassifier(algorithm='ball_tree', leaf_size=30, metric='minkowski',
                               metric_params=None, n_jobs=-1, n_neighbors=9, p=2,
                               weights='uniform')
```

Storing the model for future reference:

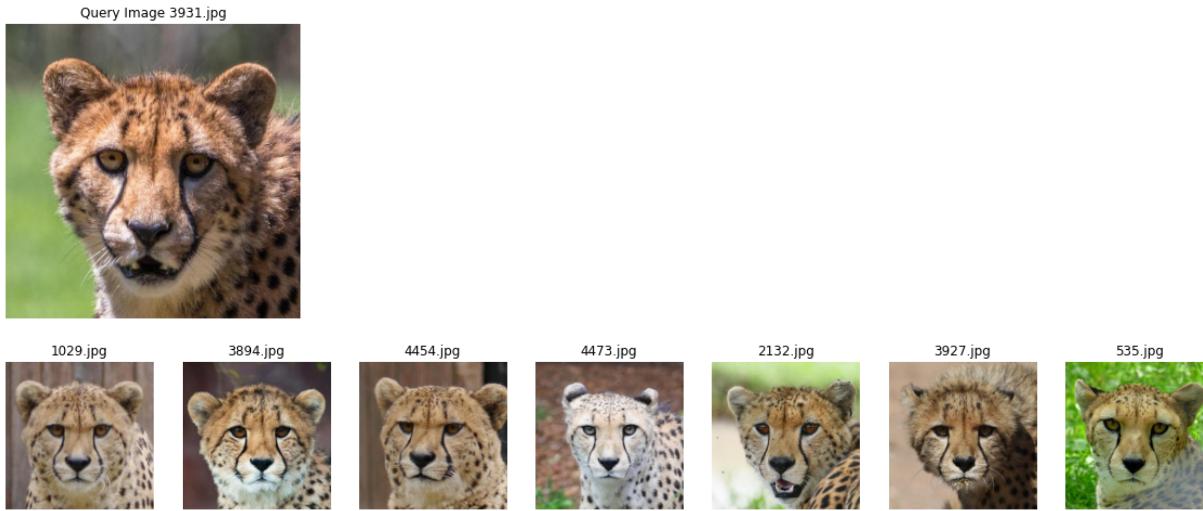
```
In [ ]: knn_file = '/content/drive/My Drive/knn_model.pkl'
joblib.dump(knn,knn_file)
```

```
In [ ]: def results_(query,result):
    """
    Plotting the N similar images from the dataset with query image.
    Arguments:
    query - (string) - filename of the query image
    result - (list) - filenames of similar images
    """
    def read(img):
        image = cv2.imread('/content/dataset/'+img)
        image = cv2.cvtColor(image,cv2.COLOR_BGR2RGB)
        return image
    plt.figure(figsize=(10,5))
    if type(query)!=type(30):
        plot_(query,"",1,1,1,"Query Image","", "",True)
    else:
        plot_(read(files[query]),"","",1,1,1,"Query Image "+files[query],"","", "",True)
    plt.show()
    plt.figure(figsize=(20,5))
    for iter,i in enumerate(result):
        plot_(read(files[i]),"","",1,len(result),iter+1,files[i],"","", "",True)
    plt.show()
```

Image Similarity Model

Making predictions:

```
In [ ]: num = 10 #datapoint
res = knn.kneighbors(data[num].reshape(1,-1),return_distance=True,n_neighbors=8)
results_((num,list(res[1][0])[1:]))
```



```
In [ ]: def predictions(label,N=8,isurl=False):
    """
    Making predictions for the query images and returns N similar images from the dataset.
    We can either pass filename or the url for the image.
    Arguments:
    label - (string) - file name of the query image.
    N - (int) - Number of images to be returned
    isurl - (string) - if query image is from google is set to True else False(By default = False)
    """
    if isurl:
        img = io.imread(label)
        img = cv2.resize(img,(224,224))
    else:
        img_path = '/content/dataset/'+label
        img = image.load_img(img_path, target_size=(224,224))
        img_data = image.img_to_array(img)
        img_data = np.expand_dims(img_data,axis=0)
        img_data = preprocess_input(img_data)
        feature = K.function([model.layers[0].input],[model.layers[12].output])
        feature = np.array(feature).flatten().reshape(1,-1)
        res = knn.kneighbors(feature.reshape(1,-1),return_distance=True,n_neighbors=N)
        results_((img,list(res[1][0])[1:]))
```

```
In [ ]: query_path = '2427.jpg'  
predictions(query_path)
```

Query Image



3639.jpg

2575.jpg

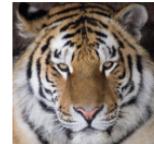
4337.jpg

3652.jpg

3373.jpg

1502.jpg

116.jpg



```
In [ ]: query_path = '4160.jpg'  
predictions(query_path)
```

Query Image



4647.jpg

305.jpg

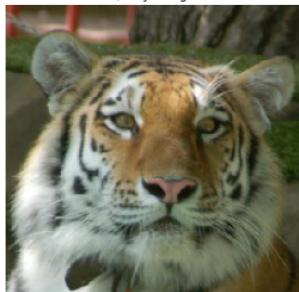
2204.jpg

3124.jpg



```
In [ ]: query_path = '1127.jpg'  
predictions(query_path)
```

Query Image



3284.jpg

1694.jpg

554.jpg

116.jpg

4083.jpg

3664.jpg

132.jpg



```
In [ ]: query_path = '1379.jpg'  
predictions(query_path,6)
```

Query Image



578.jpg



2858.jpg



4128.jpg



3985.jpg

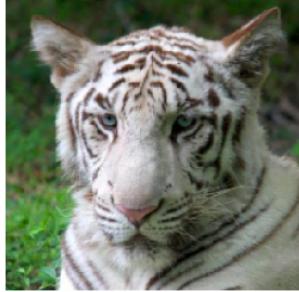


4528.jpg



```
In [ ]: query_path = '608.jpg'  
predictions(query_path)
```

Query Image



1670.jpg



108.jpg



1738.jpg



3476.jpg



3727.jpg



147.jpg



4178.jpg



```
In [ ]: query_path = '2550.jpg'  
predictions(query_path)
```



135.jpg

1515.jpg

747.jpg

262.jpg

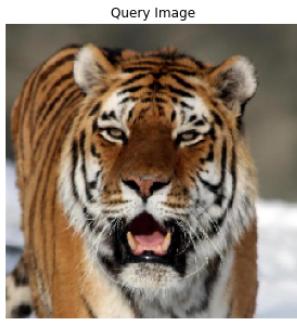
1966.jpg

4584.jpg

3985.jpg



```
In [ ]: query_path = '167.jpg'  
predictions(query_path)
```



274.jpg

680.jpg

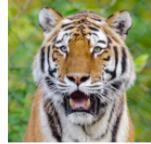
2727.jpg

3484.jpg

1097.jpg

3222.jpg

3362.jpg



```
In [ ]: query_path = '543.jpg'  
predictions(query_path,7)
```



1570.jpg

3084.jpg

3262.jpg

3058.jpg

1878.jpg

2679.jpg



Testing with google images:

```
In [ ]: import imageio as io  
query_path = 'https://tse4.mm.bing.net/th?id=OIP.NIMP0bTfhF3898t_ZYLB8QHaE8&pid=Api&P=0&w=248&h=166'  
predictions(query_path,4,isurl=True)
```

Query Image



3432.jpg



3899.jpg



3485.jpg



```
In [ ]: query_path = 'https://tse4.mm.bing.net/th?id=OIP.dr7YKzR28BFJerLKJc4cLgHaE7&pid=Api&P=0&w=242&h=162'  
predictions(query_path,5,isurl=True)
```

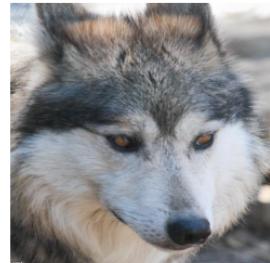
Query Image



4439.jpg



1846.jpg



1337.jpg



1017.jpg



```
In [ ]: query_path = 'https://tse2.mm.bing.net/th?id=OIP.EkSb6SADo3WRj2WFQWTBvgHaFj&pid=Api&P=0&w=216&h=163'  
predictions(query_path,6,isurl=True)
```

Query Image



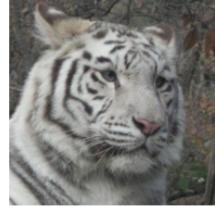
1738.jpg



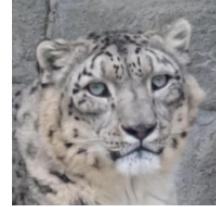
4178.jpg



3955.jpg



2277.jpg



978.jpg



```
In [ ]:
```